LSTM diagram

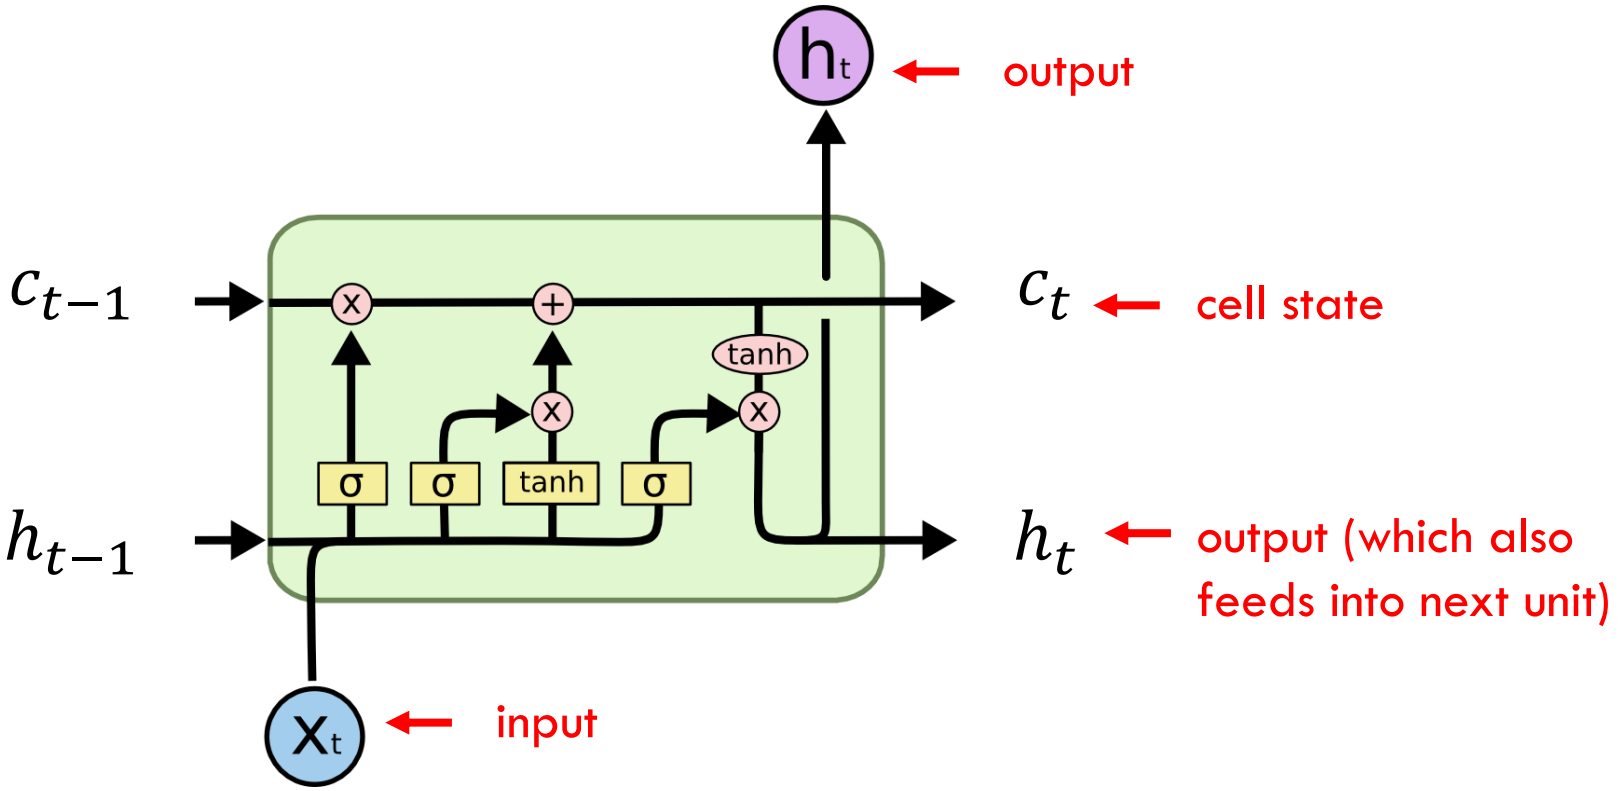

LSTM diagram

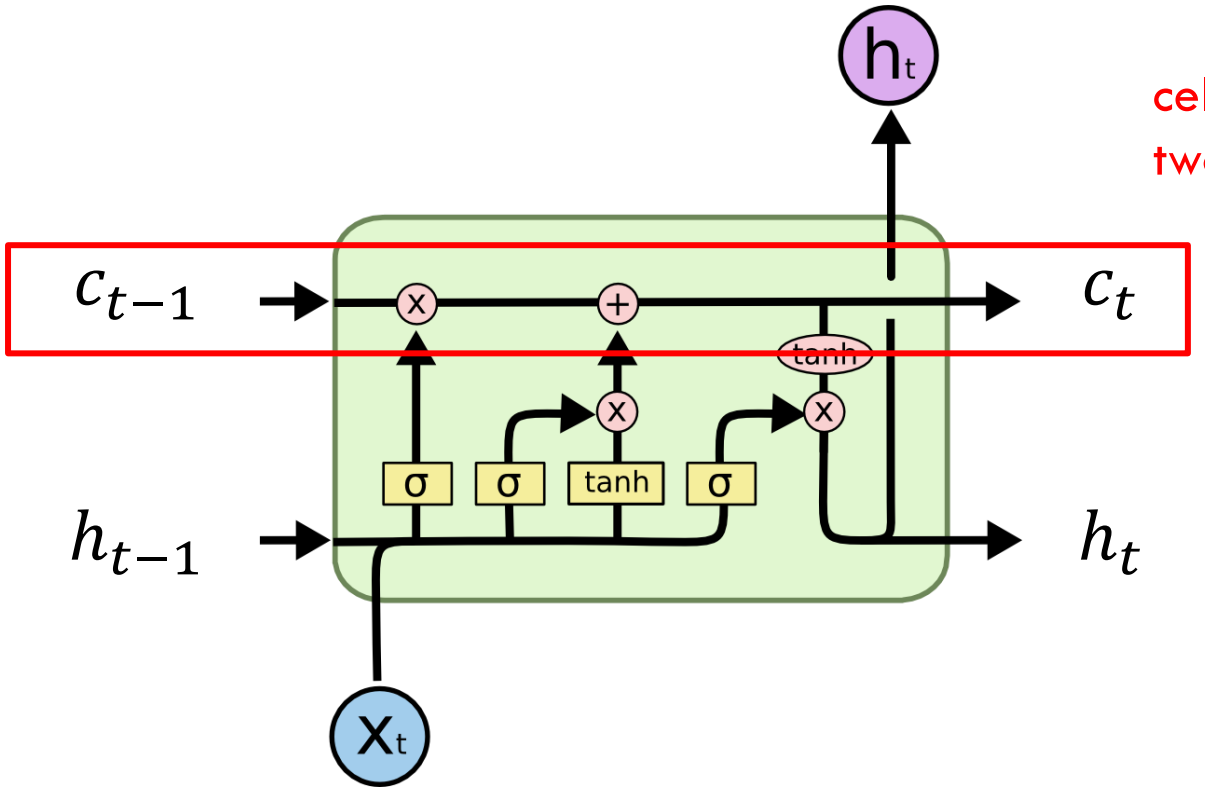

cell state gets updated in two stages

LSTM diagram

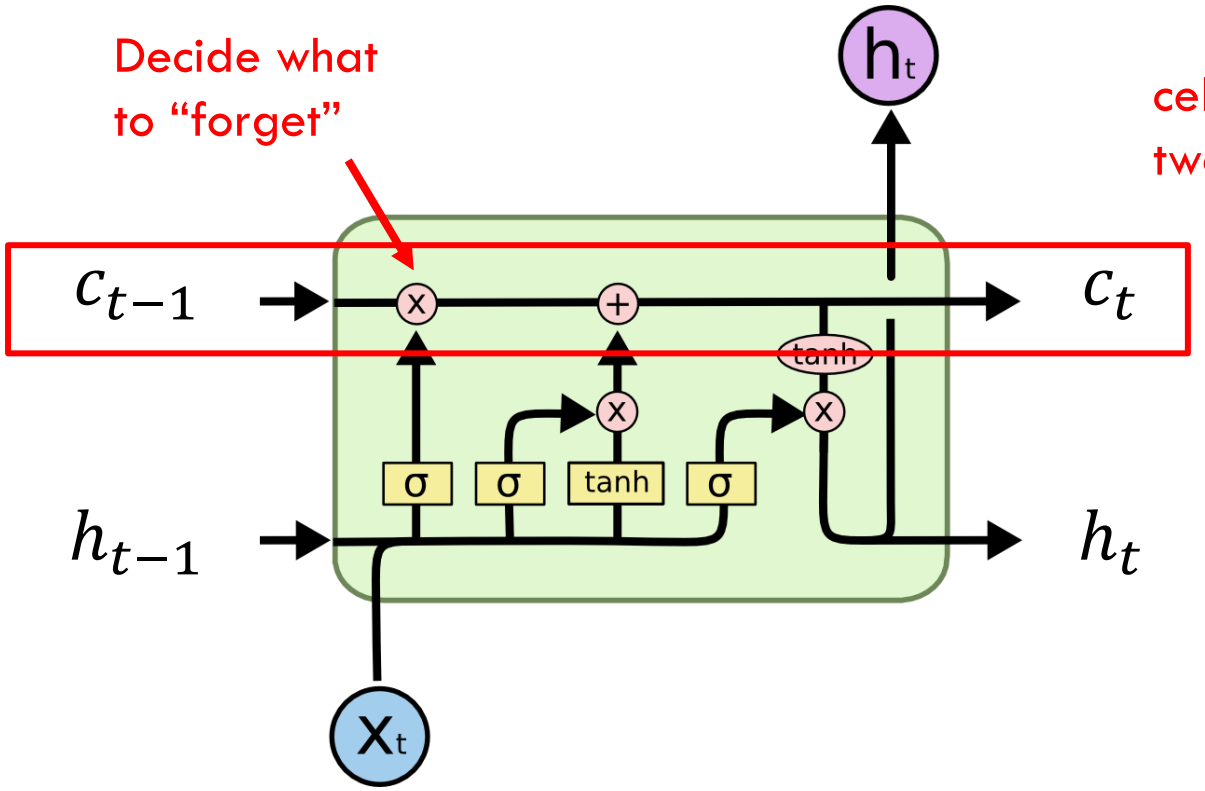

LSTM diagram

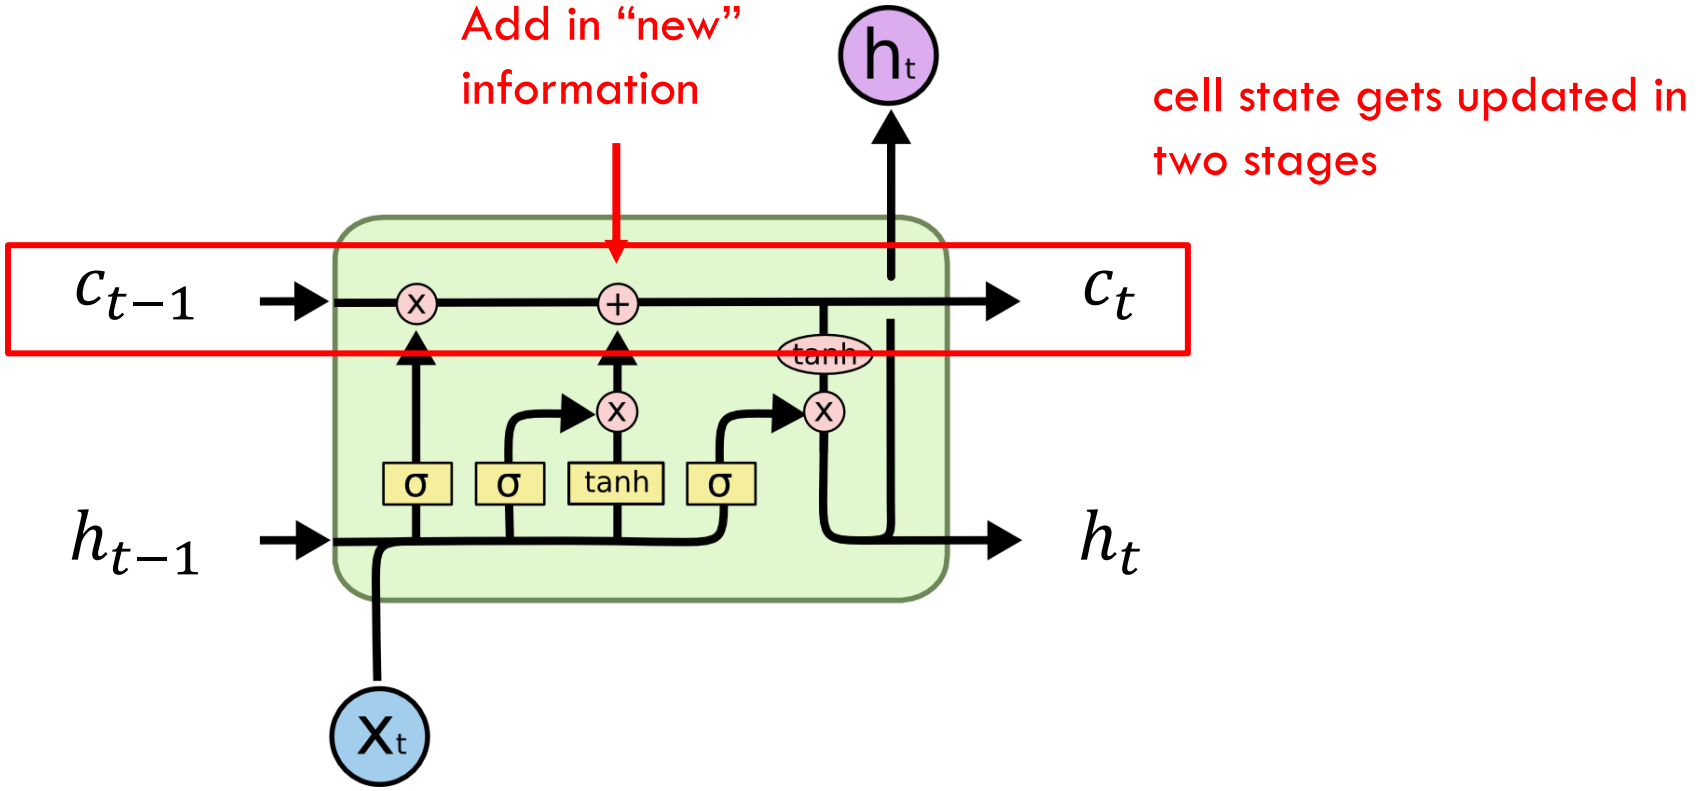

LSTM diagram

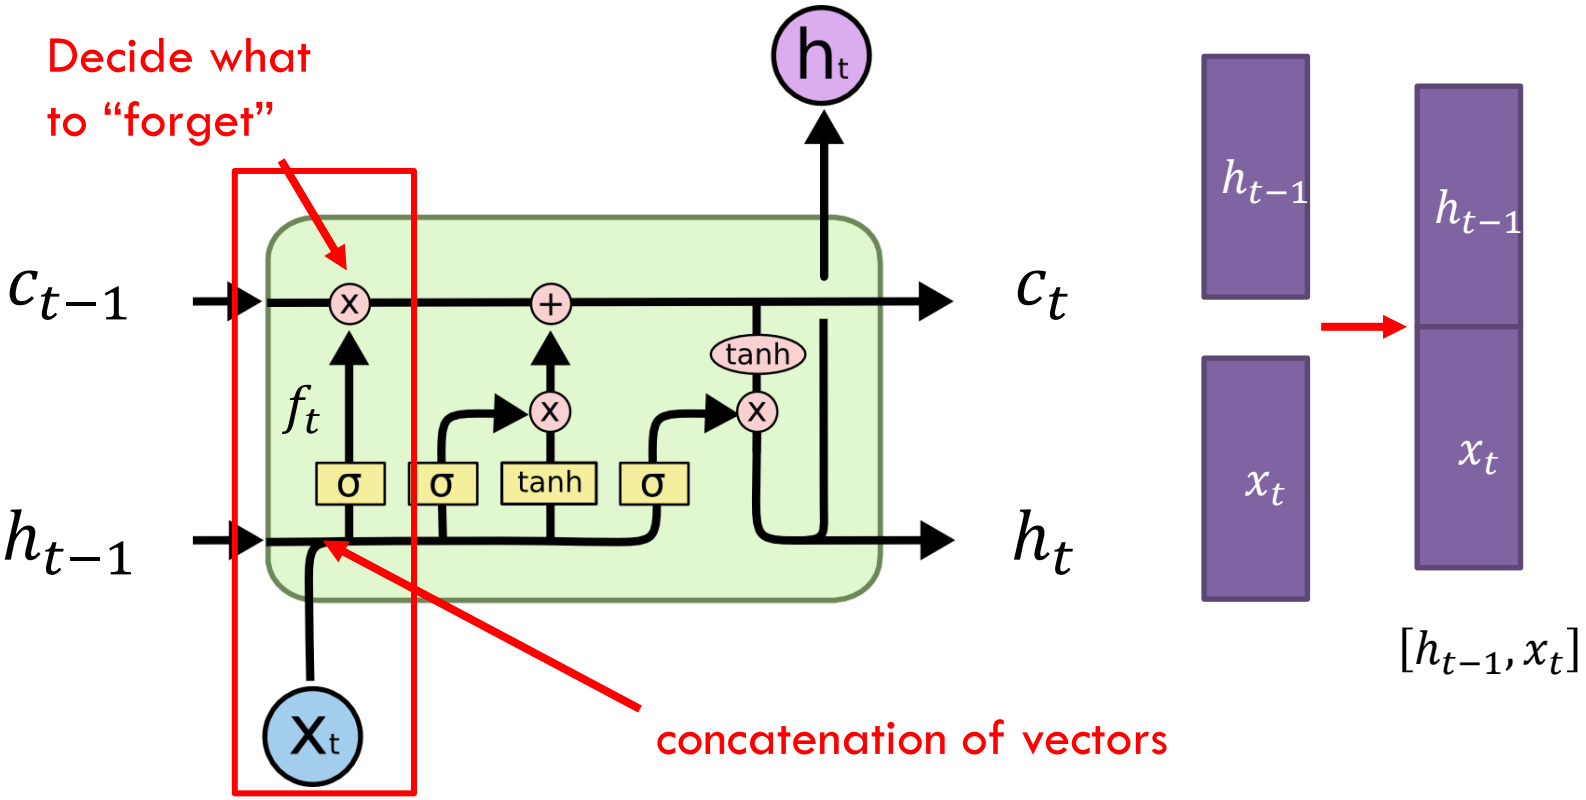

LSTM diagram

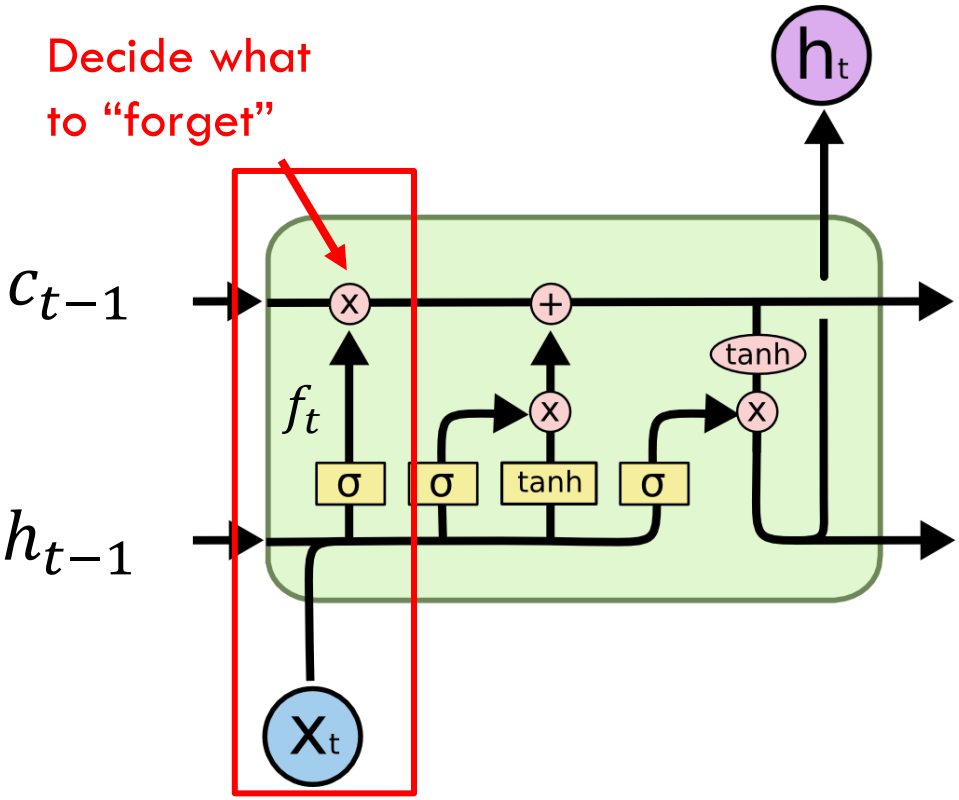

$$f_t = \sigma\left(W_f[h_{t-1}, x_t] + b_f\right)$$

based on
previous output
and current
input

LSTM diagram

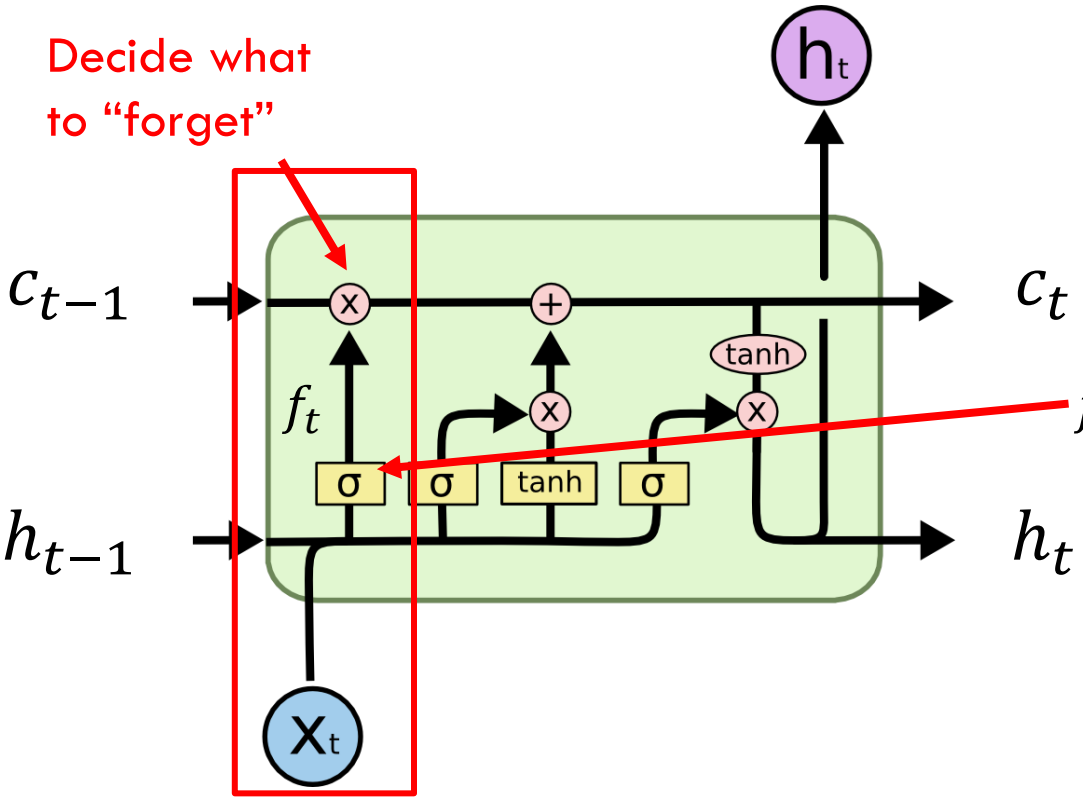

$$f_t = \sigma\left(W_f[h_{t-1}, x_t] + b_f\right)$$

based on
previous output
and current
input

LSTM diagram

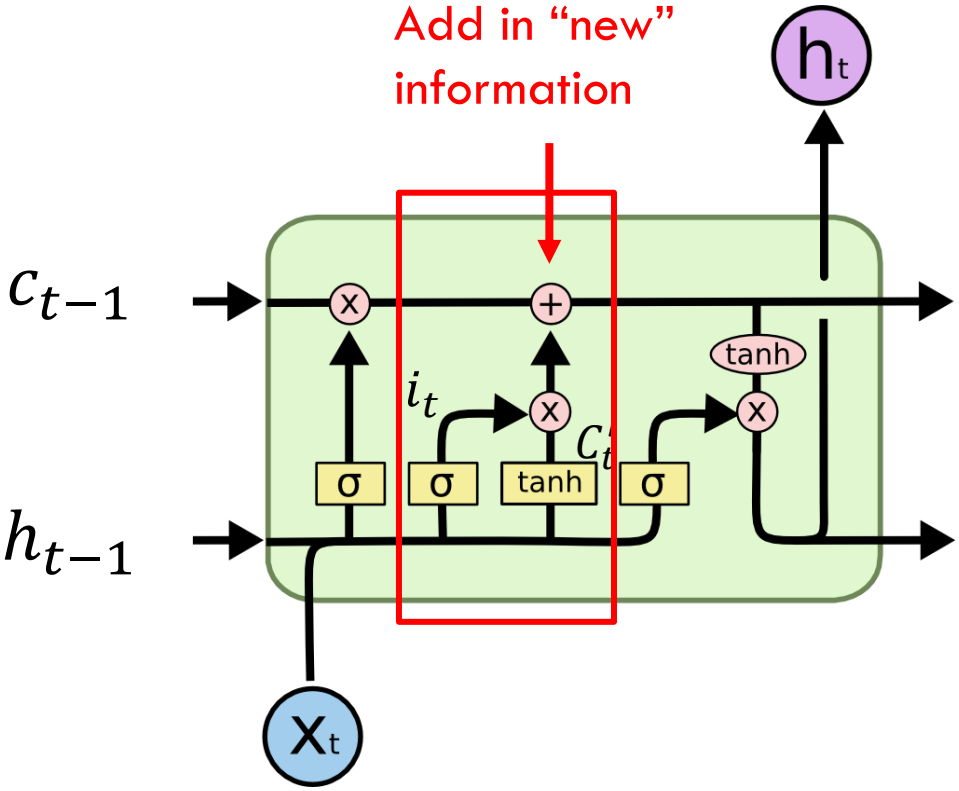

$$i_t = \sigma(W_i[h_{t-1}, x_t] + b_i)$$

$$C'_t = tanh(W_C[h_{t-1}, x_t] + b_C)$$

LSTM diagram

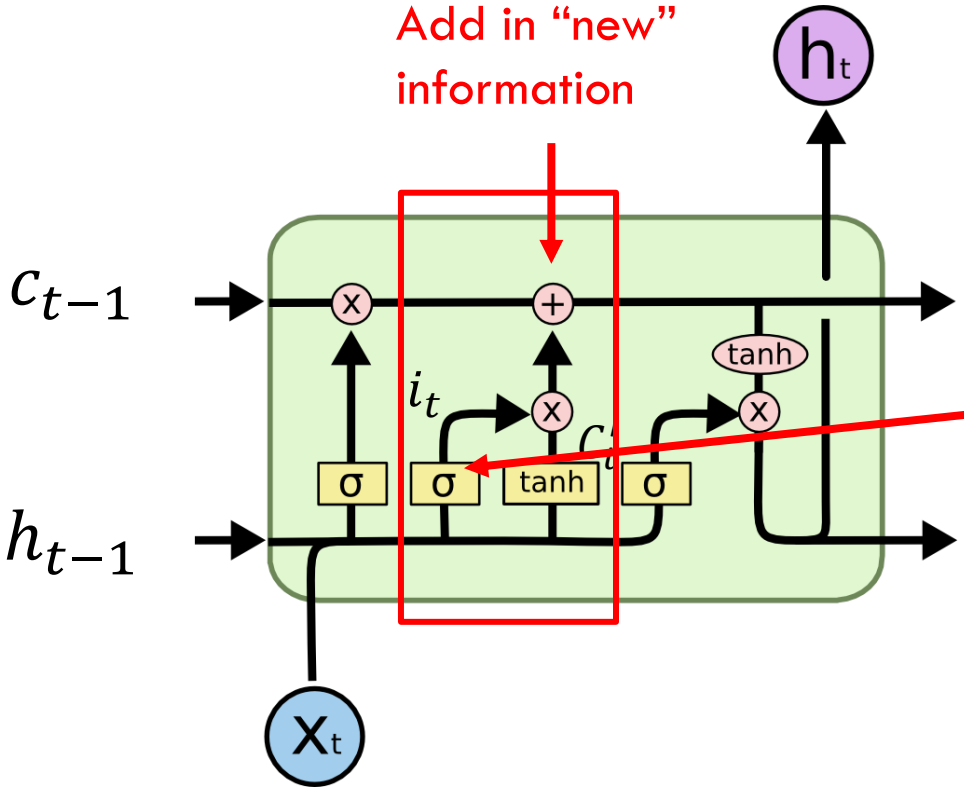

$$i_t = \sigma(W_i[h_{t-1}, x_t] + b_i)$$

$$C'_t = tanh(W_C[h_{t-1}, x_t] + b_C)$$

LSTM diagram

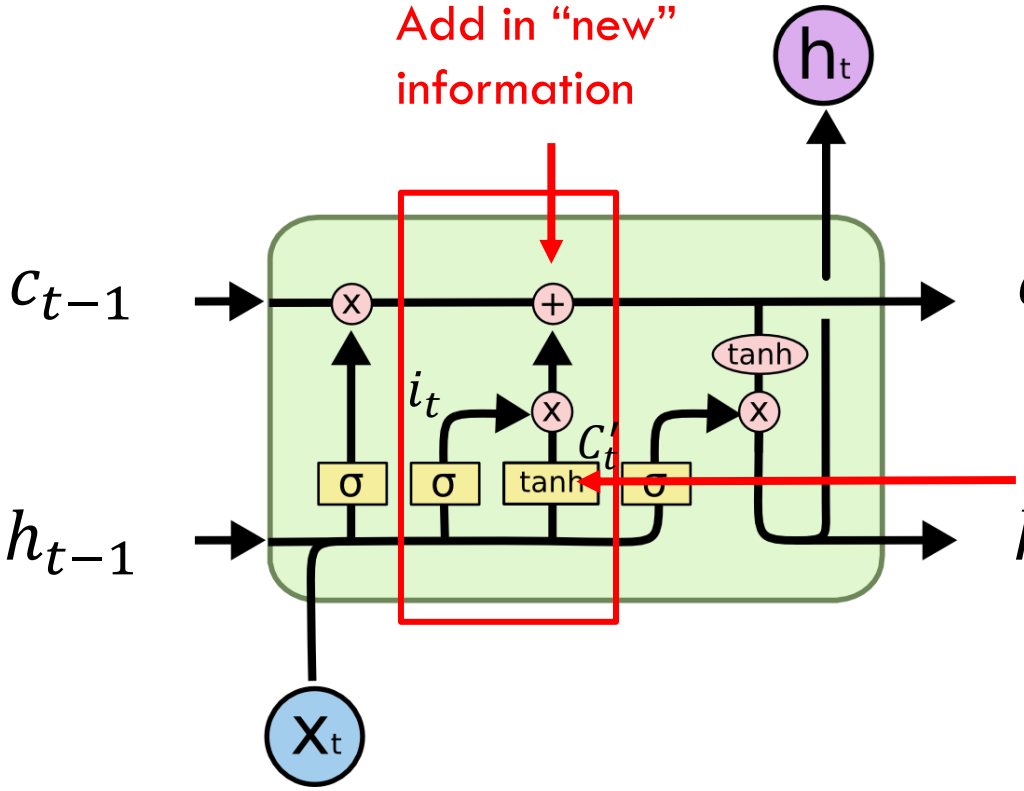

$$i_t = \sigma(W_i[h_{t-1}, x_t] + b_i)$$

$$C'_t = tanh(W_C[h_{t-1}, x_t] + b_C)$$

LSTM diagram

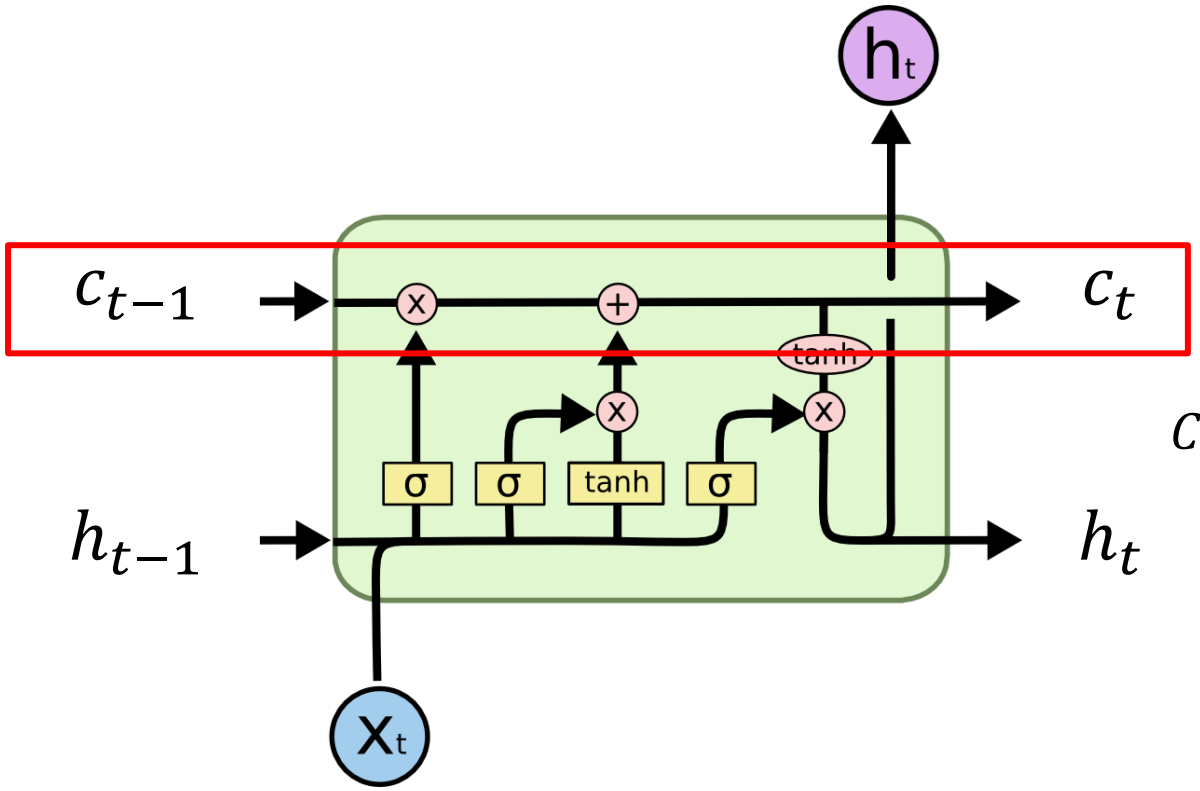

Note: '*' represents element-wise multiplication

$$C_t = f_i * C_{t-1} + i_t * C'_t$$

forget the old (or not)

add the new (or not)

LSTM diagram

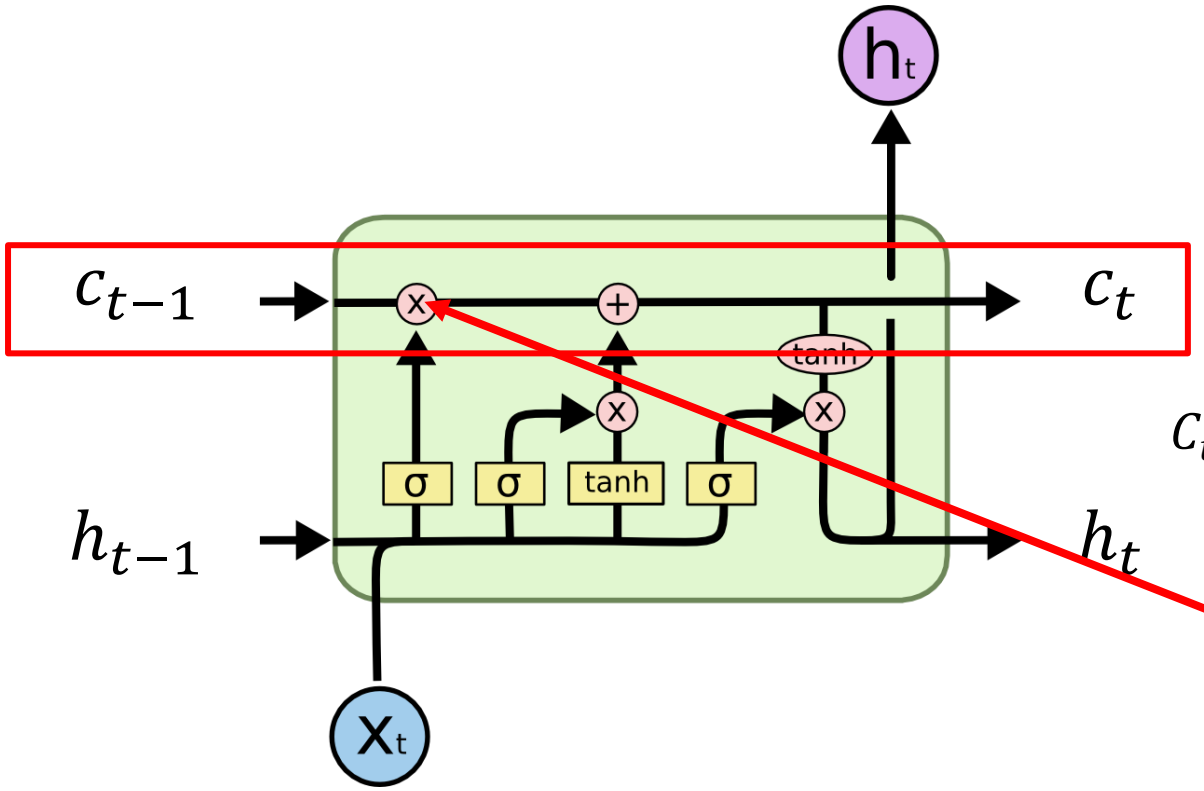

Note: '*' represents element-wise multiplication

$$C_t = f_i * C_{t-1} + i_t * C'_t$$

forget the old (or not)

add the new (or not)

LSTM diagram

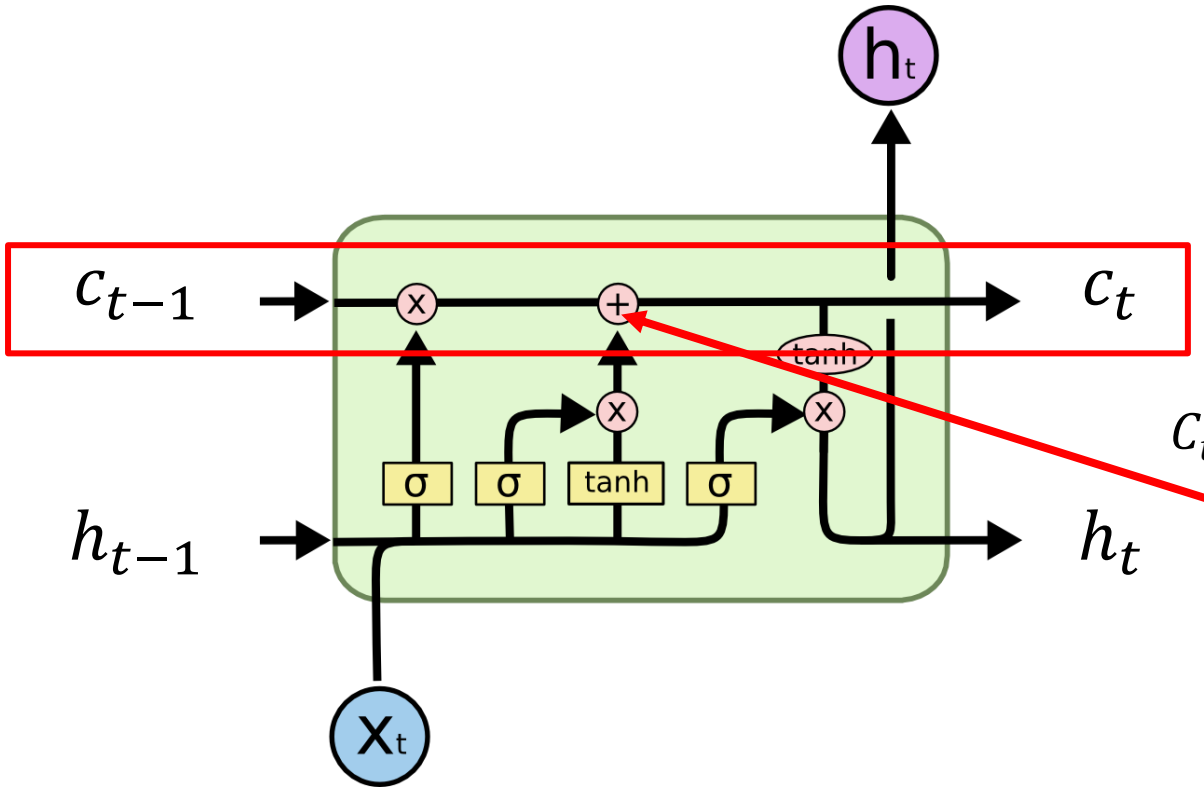

Note: '*' represents element-wise multiplication

$$C_t = f_i * C_{t-1} + i_t * C'_t$$

forget the old (or not)

add the new (or not)

LSTM diagram

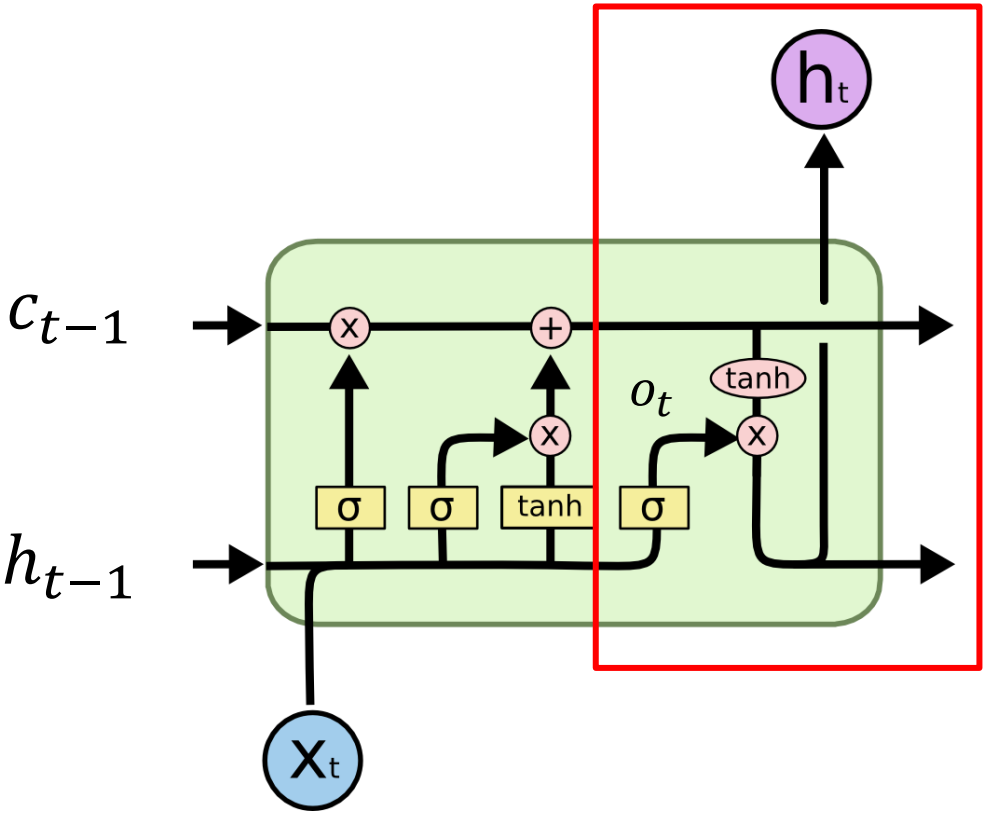

Final stage computes the output

$$o_t = \sigma(W_o[h_{t-1}, x_t] + b_o)$$

$$h_t = o_t * \tanh(C_t)$$

LSTM diagram

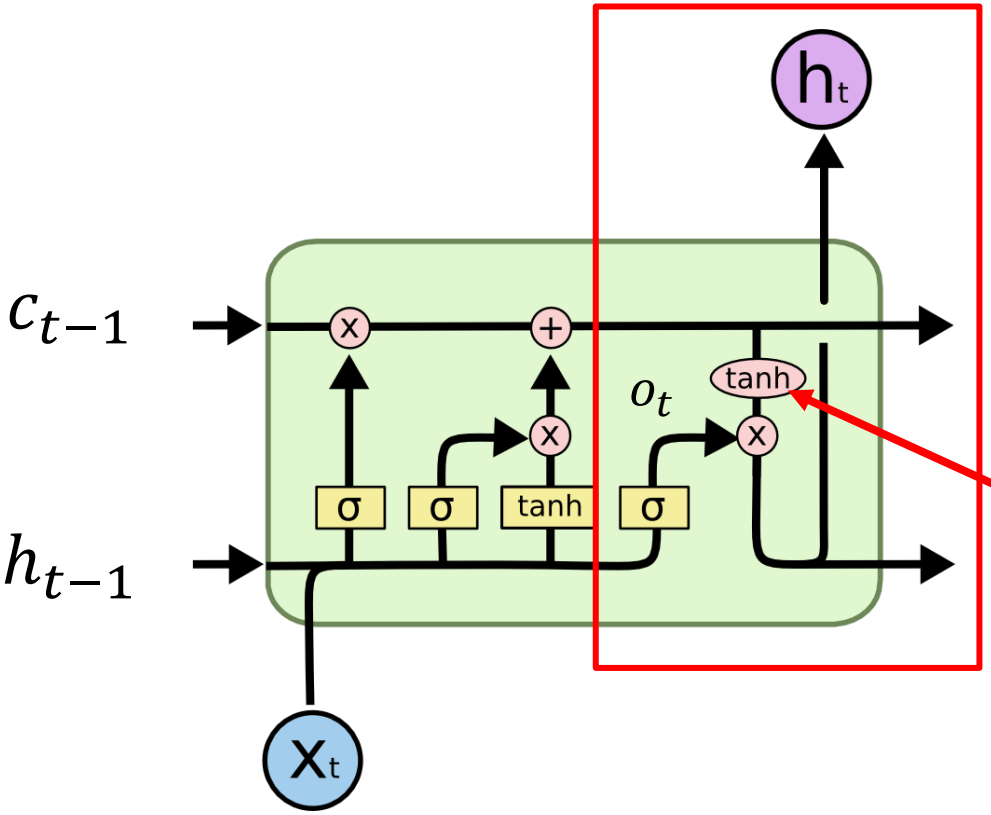

Final stage computes the output

$$o_t = \sigma(W_o[h_{t-1}, x_t] + b_o)$$

$$h_t = o_t * \tanh(C_t)$$

Note: No weights here

LSTM unrolled

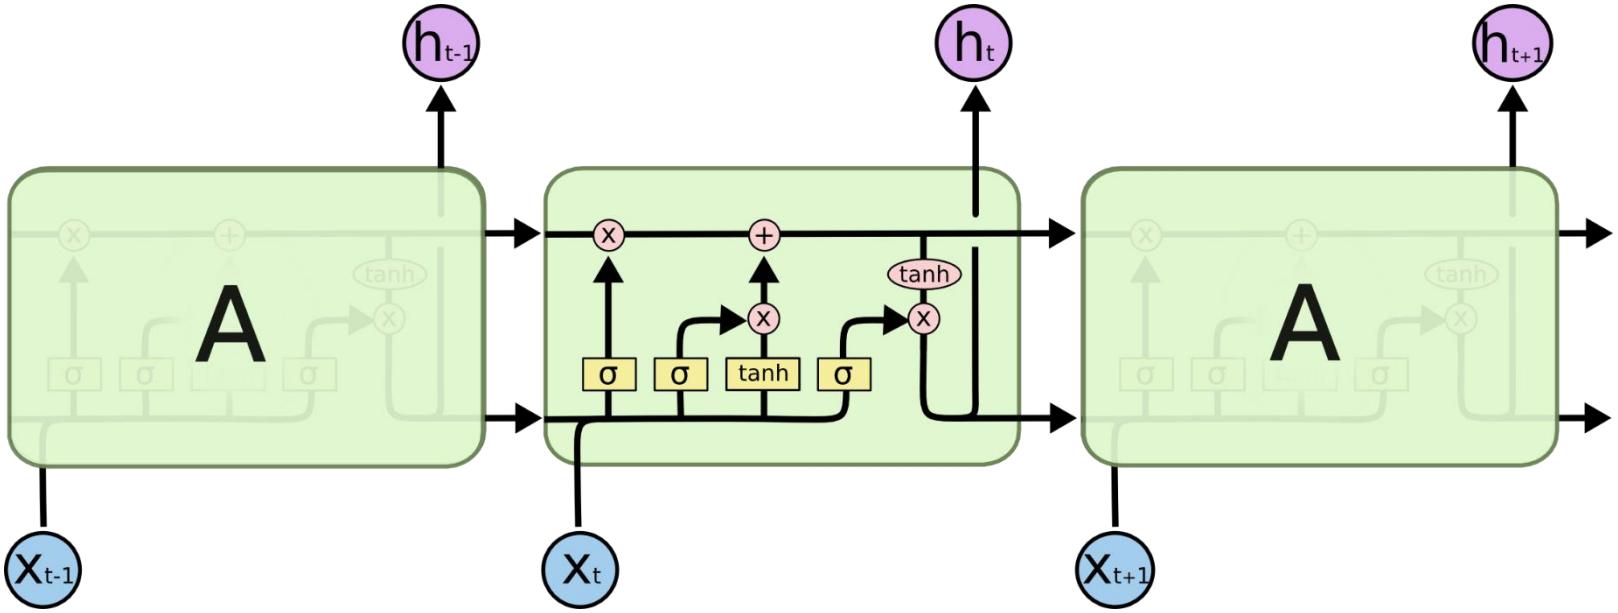

Final Points

- This is the most common version of LSTM, but there are many different “flavors”.
  - Gated Recurrent Unit (GRU)
  - Depth-Gated RNN

- LSTMs have considerably more parameters than plain RNNs.

- Most of the big performance improvements in NLP have come from LSTMs, not plain RNN.